

OpenACC/OpenMP

the Easy Ways to GPU Programming

Giovanni Tedeschi Prades


Code Coffee -25.11.2024

Outlook

- A basic overview of parallelization in CPUs and GPUs
- GPU resources at the Physics Cluster
- **First hands-on:** getting a GPU, open a parallel region, compile and run. You can code along if you want.
- The Mandelbrot set
- **Second hands-on:** parallelize a for loop
- Checking the performances

Disclaimers

- I am **not** an expert in GPU programming, I just attended a school
- Ask and interrupt, I'll do my best
- I'll be working in C



An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as

OpenMPI, MPICH...

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as

OpenMPI, MPICH...

Not the same thing

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as
OpenMPI, MPICH...

GPU

CUDA-programming

A «programming model» to directly
interact with GPUs

Available only for NVIDIA GPUs

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as
OpenMPI, MPICH...

GPU

CUDA-programming

A «programming model» to directly
interact with GPUs

Available only for NVIDIA GPUs

OpenMP and **OpenACC**

Based on compiler directives which
tell the compiler the sections of the
code to run on the GPU

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as
OpenMPI, MPICH...

GPU

CUDA-programming

A «programming model» to directly
interact with GPUs

Available only for NVIDIA GPUs

OpenMP and **OpenACC**

Based on compiler directives which
tell the compiler the sections of the
code to run on the GPU

Actually the same thing

An incomplete map of parallelization

CPU

Shared-memory parallelization

OpenMP

Implemented in most compilers

Distributed-memory parallelization

MPI (**M**essage **P**arsing **I**nterface)

Implemented in modules such as
OpenMPI, MPICH...

GPU

CUDA-programming

A «programming model» to directly
interact with GPUs

Available only for NVIDIA GPUs

OpenMP and **OpenACC**

Based on compiler directives which
tell the compiler the sections of the
code to run on the GPU

Kokkos (and friends)

High-level programming frameworks to
make a code highly portable, including
GPUs

An incomplete map of parallelization

CUDA

OpenMP/OpenACC

Kokkos



An incomplete map of parallelization

Take this with a grain of salt

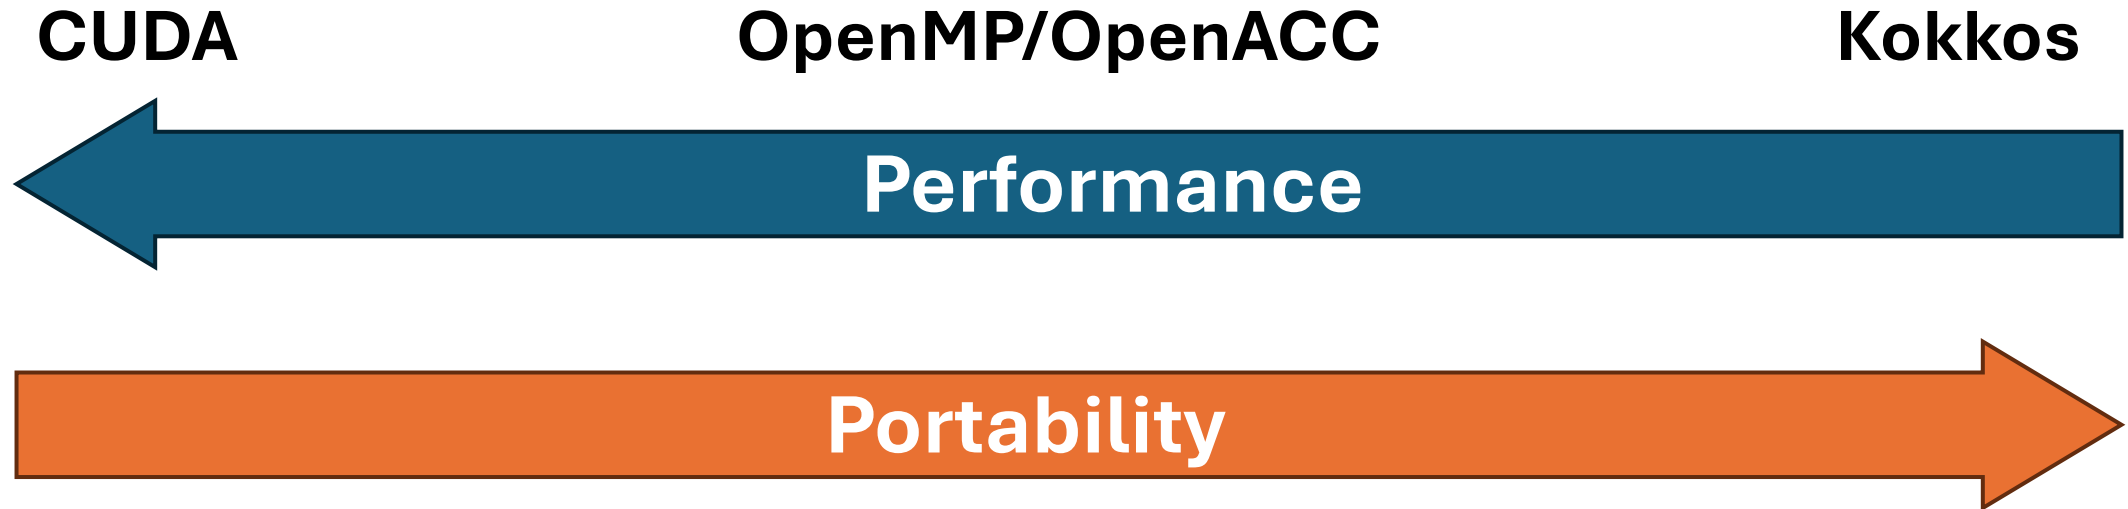
CUDA

OpenMP/OpenACC

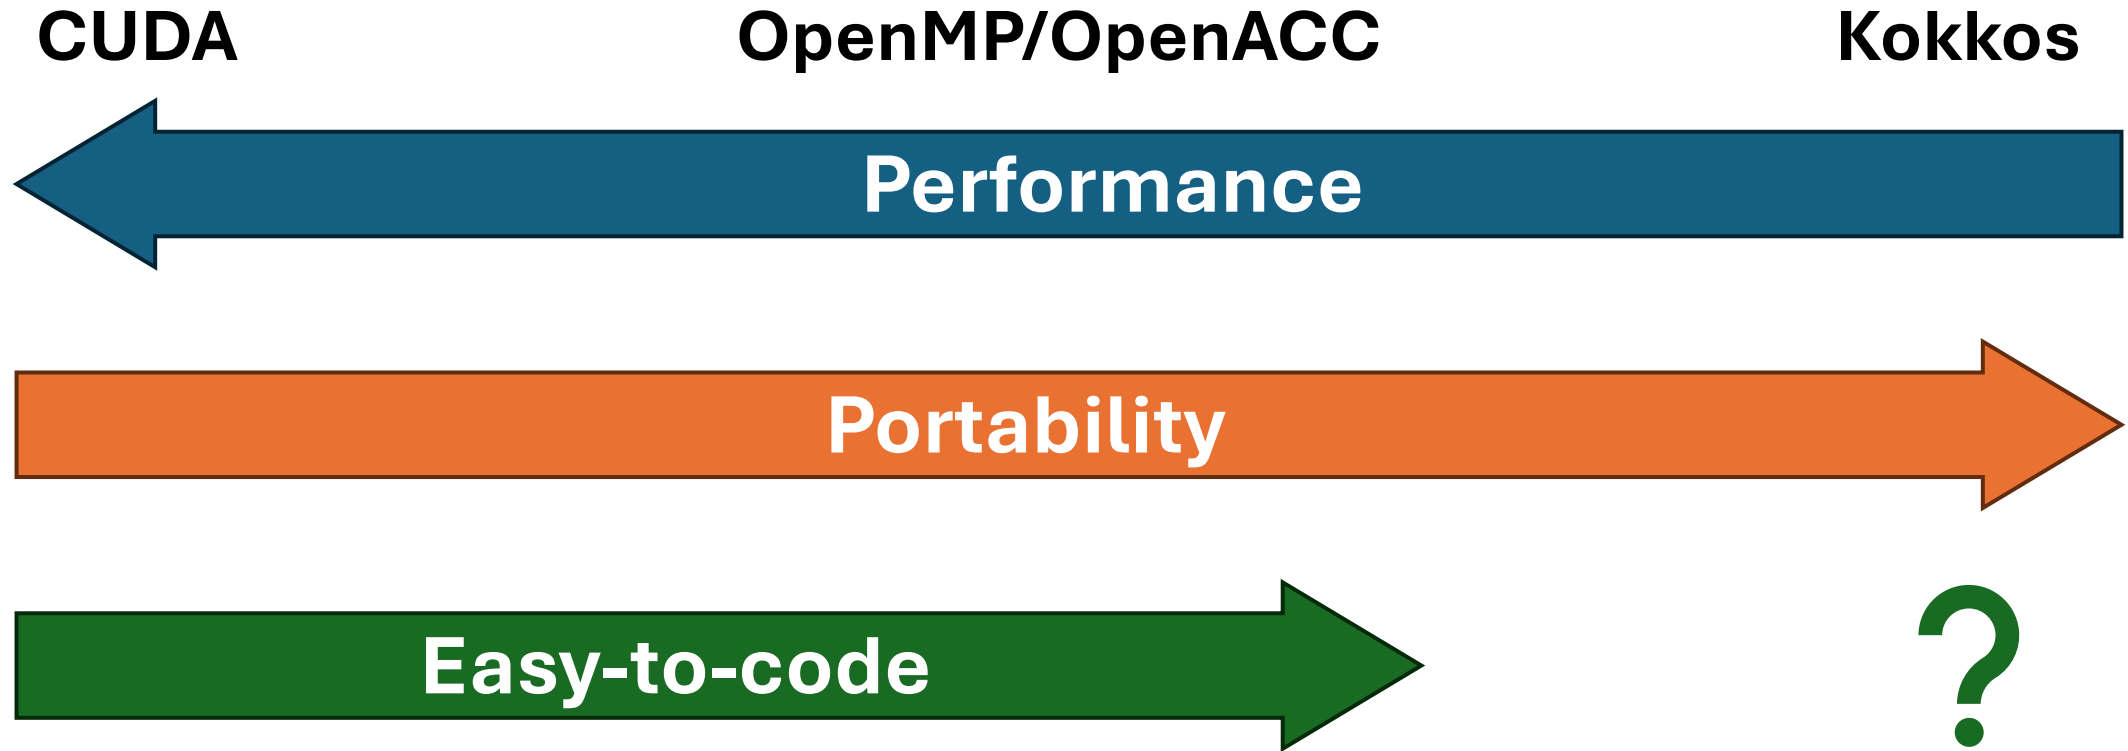
Kokkos



An incomplete map of parallelization




An incomplete map of parallelization



GPUs at the Physics Cluster

- 121 total GPUs available at the cluster
 - 28x P2000
 - 4x P5000
 - 8x V100
 - 7x H100
 - 66x A40
 - 2x A100
 - 4x RTX2080Ti
 - 2x TitanXP
- All NVIDIA GPUs

GPUs at the Physics Cluster

- 121 total GPUs available at the cluster
 - 28x P2000
 - 4x P5000
 - 8x V100
 - 7x H100
 - 66x A40  49 of which available at the CIP nodes
 - 2x A100
 - 4x RTX2080Ti
 - 2x TitanXP
- All NVIDIA GPUs

GPUs at the Physics Cluster

- 121 total GPUs available at the cluster
 - 28x P2000
 - 4x P5000
 - 8x V100
 - 7x H100
 - 66x A40
 - 2x A100
 - 4x RTX2080Ti
 - 2x TitanXP
 - All NVIDIA GPUs
-] Not compatible with OpenMP

To the terminal!

- **How to start an interactive job with a GPU**

Simple version: `intjob --gres=gpu:1`

To specify which GPU you want (in this case an A40, available only at the `inter` partition):

```
intjob --gres=gpu:a40:1 --partition=inter
```

- **Compilers and commands**

Compile with OpenACC: `nvc -acc test.c`

Compile with OpenMC: `nvc -mp=gpu test.c`

The `nvc` compiler is available in the `nvhpc` module at the cluster

To the terminal!

- **Open a parallel region with OpenACC and OpenMP**

OpenACC: `#pragma acc parallel`

OpenMP: `#pragma omp target`


- **Compile, run and check you are actually using a GPU**

Use the command `nvidia-smi` to check the status of the GPU.

Useful to put it in a watch `-n 0.1 nvidia-smi` to see it changing live.


The Mandelbrot set

- Defined as the set of points z_0 in the complex plane for which the sequence $z_{n+1} = z_n^2 + z_0$ converges.
- Computationally demanding at the edges, where a large number of iterations are required
- Embarassingly parallel: each point can be evaluated independently to the others
- Double for loop to check all the points inside a certain domain



The Mandelbrot set

- We will focus somewhere here



To the terminal!

- **How to parallelize for loops with GPUs**

OpenACC: `#pragma acc parallel loop`

OpenMP: `#pragma omp target parallel for`

- **The basics of data exchange**


OpenACC: `copyin(<var_name>) copyout(<var_name>)`

OpenMP: `map(to:<var_name>) map(from:<var_name>)`


Needed for OpenMP, otherwise it will crash!

- **What kind of speedup is possible?**

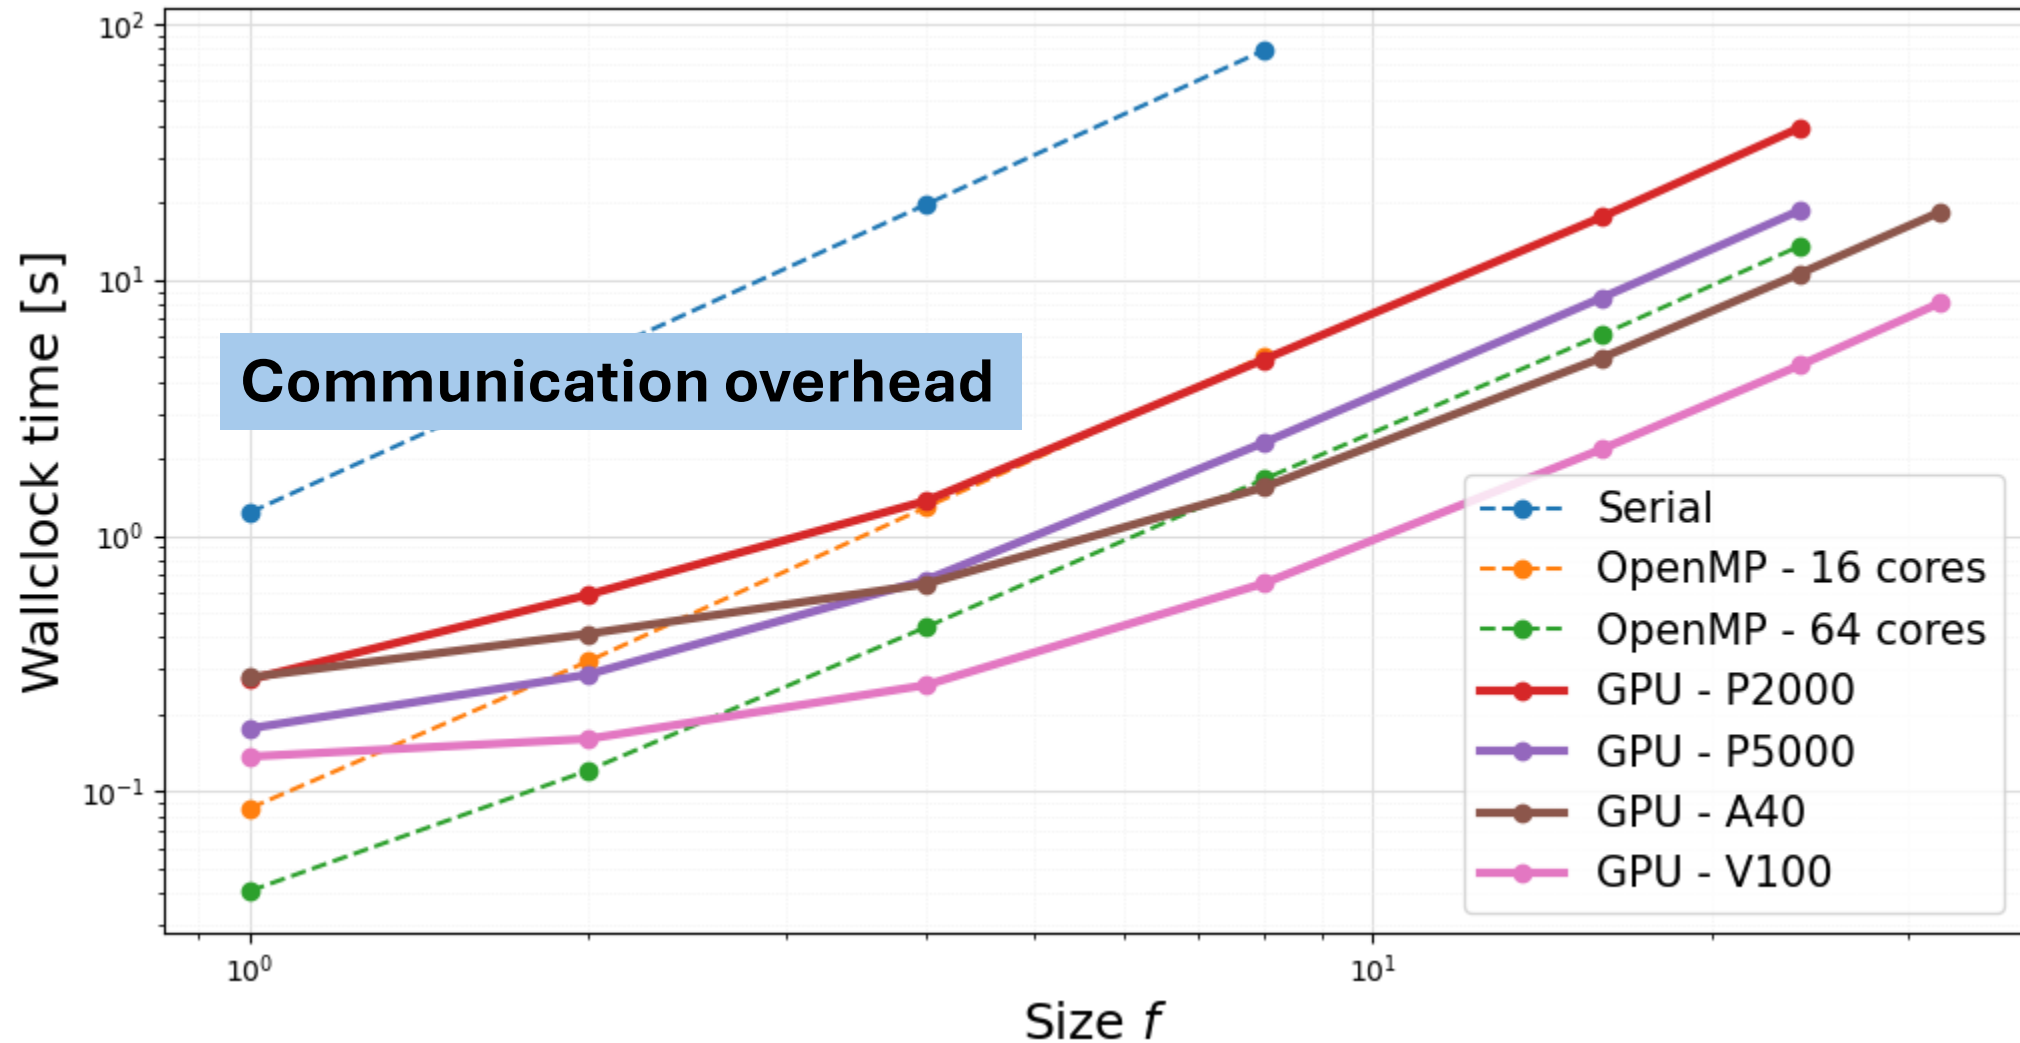
Performance study



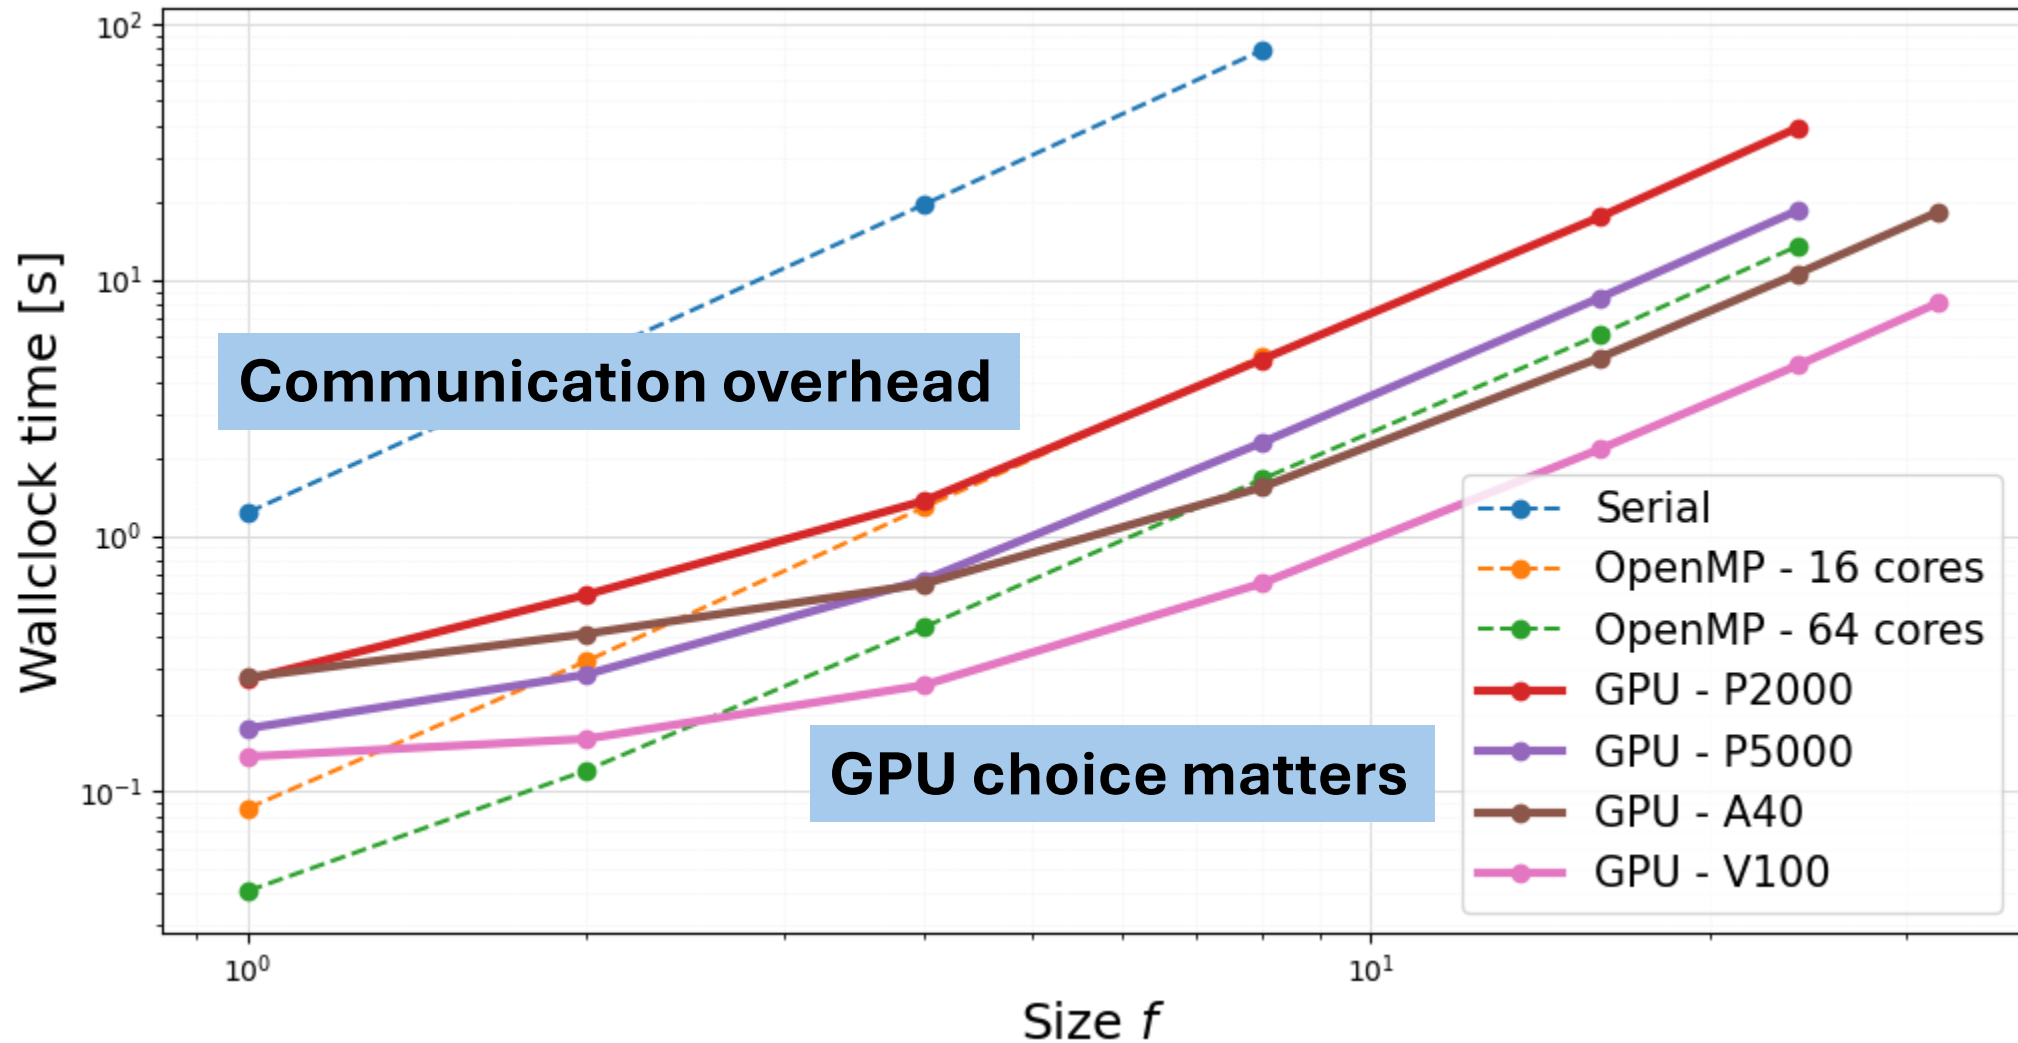
Performance study



Performance study



Performance study



#pragma summary

OpenACC

```
#pragma acc parallel
```

```
#pragma acc parallel loop  
copyin(<var_name>)  
copyout(<var_name>)
```

OpenMP

```
#pragma omp target
```

```
#pragma omp target parallel for  
map(to:<var_name>)  
map(from:<var_name>)  
map(tofrom:<var_name>)
```

Conclusions

- OpenMP and OpenACC offer an easy way into GPU programming
- Most of the heavy work is done by the compiler
- Compilers matter: nvc is the best choice for NVIDIA GPUs
- Many #pragma directives offered, syntax differs between the two
- Mind the communication overhead
- The choice of which GPU to use matters
- A useful resource: [OpenACC programming and Best Practices Guide](#)